Project Summary

Sample-Efficiency of Quantum Fidelity Kernels
for Molecular Property Classification:
A Learning-Curve Benchmark Against Classical Baselines

Supaporn Klabklaydee
Institute of Science Tokyo

A negative / null-result study with multi-seed, scaffold-split learning curves.

Core Question

Does a quantum fidelity kernel (simulated on a classical computer) show a real *sample-efficiency* advantage over classical machine-learning methods when labeled molecular data is scarce—the low-data regime that matters most in drug discovery and toxicology, where labels are expensive?

Method

The study benchmarks a PennyLane simulator-based quantum fidelity kernel SVM (8-qubit "Circuit B," using angle encoding + ring CNOT entanglement on PCA-8 molecular features) against three classical baselines—Random Forest, RBF-SVM, and AttentiveFP graph neural network— across four MoleculeNet binary classification tasks:

- **BBBP** — blood–brain barrier permeability
- **BACE** — enzyme inhibition
- **HIV** — replication inhibition
- **Tox21 NR-AR** — toxicity

Key design rigor. Training sizes were swept from $N = 50$ up to the full dataset (7 sizes), with Bemis–Murcko scaffold splitting (to prevent structural leakage between train and test) and 10 random seeds—**280 total model–condition pairs**. A pre-registered hypothesis was tested: that the quantum kernel would show a statistically significant advantage over the GNN specifically at $N < 500$ (paired Wilcoxon, $p < 0.05$).

Key Results

Dataset	Quantum vs. GNN advantage	Quantum vs. best classical model
BBBP	Significant at $N = 300$, $N = 500$	Never exceeded RF / RBF-SVM
BACE	No significant advantage	Quantum near chance-level ($\sim$0.50 AUC) throughout
HIV	Significant at $N = 50$–200	Never exceeded RF / RBF-SVM
Tox21	Mean crossover at $N = 50$, not significant ($p = 0.23$)	Not robust; RBF-SVM later dominated

The quantum kernel's AUC stayed flat in a narrow **0.46–0.53** band on three of four tasks—essentially chance-level—while Random Forest and RBF-SVM generally dominated overall, with gaps of **0.12–0.34 AUC** points at larger N.

Why the "Significant" Wins Are Not What They Look Like

The paper is unusually candid: the quantum-over-GNN significant results largely reflect **GNN underfitting at small** N, not quantum kernel strength. On HIV at $N = 50$, the GNN scored only **0.340** while Random Forest already hit **0.552**—the quantum kernel's **0.491** beat the GNN but still trailed the simpler classical baselines.

Explanation Offered

The authors attribute the quantum kernel's plateau to **expressivity limits**: with 8 qubits, the fidelity kernel's rank is bounded by a 256-dimensional Hilbert space and constrained by the fixed ring-CNOT circuit structure, unlike the RBF kernel or Morgan fingerprints, which operate in much higher-dimensional, more flexible spaces. This points to a *capacity-saturated* regime rather than overfitting.

Practical Trade-offs

Computing the quantum kernel's Gram matrix required $O(N^2)$ circuit evaluations, taking **hours per dataset** on a simulator, versus **seconds to minutes** for classical Random Forest and RBF-SVM on the same hardware— a substantial computational cost for no net performance benefit.

Conclusion

This is presented as the first systematic scaffold-split learning-curve comparison of a quantum fidelity kernel against modern classical baselines with proper multi-seed statistics.

Bottom line: the quantum kernel is not globally competitive with the best classical models on any of the four benchmarks, though it does show narrow, likely GNN-underfitting-driven advantages in a few low-data corners (BBBP, HIV). The authors are transparent about limitations—no Bonferroni correction across the many comparisons, simulator-only evaluation, and fixed (non-trainable) quantum circuit design— and frame the exploratory "wins" as evidence of GNN fragility at low N rather than proof of quantum advantage.

Takeaway. This is a negative / null-result paper with strong methodological rigor—valuable as a reproducible baseline for the field precisely because it resists overselling a fashionable "quantum advantage" narrative.